
Announcements

- ❖ Mid-term exam: June 22, 4pm-5:40pm
 - ❖ Open book, open notes
 - ❖ No communication
- ❖ HW5 on CSP
 - ❖ Released today
 - ❖ Due June 24 at 11:59pm
- ❖ P3 on MDP and RL
 - ❖ Early release
 - ❖ Due July 3 at 11:59pm

Ve492: Introduction to Artificial Intelligence

Mid-term Review



Paul Weng

UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

What have we learned so far?

- ❖ Search and planning

- ❖ Define a state space, goal test; Find path from start to goal

- ❖ Game trees

- ❖ Define utilities; Find path from start that maximizes utility

- ❖ Decision theory and game theory

- ❖ Foundation for MEU; Basic concepts in game theory

- ❖ MDPs

- ❖ Define rewards, utility = (discounted) sum of rewards
 - ❖ Find policy that maximizes utility

- ❖ Reinforcement learning

- ❖ Just like MDPs, only T and / or R are not known in advance



- ❖ Constraint satisfaction

- ❖ Find solution that satisfies constraints; Not just for finding a sequential plan



High-Level Framework

- ❖ How to build AI system?



Search

- ❖ Environment: single-agent, fully-observable state, deterministic transition, sequential, model known
 - ❖ Search problem
 - ❖ States, transition model, goal test, initial state
 - ❖ Search tree
 - ❖ Algorithms
 - ❖ Uninformed search
 - ❖ BFS, DFS, UCS
 - ❖ Informed search
 - ❖ Greedy search, A*
 - ❖ Properties
 - ❖ Complete, optimal
 - ❖ Space and computational complexities
- graph
- tree / search graph
- heuristic fcn.

Search in Games

- ❖ Environment: multi-agent, fully-observable state, deterministic or stochastic transition, turn-taking , model known
- ❖ Multi-agent search problems as games
 - ❖ States, players, transition model, terminal test/ values, initial state
 - ❖ Game tree
- ❖ Algorithm for adversarial agent (zero-sum game)
 - ❖ Minimax search algorithm
 - ❖ Alpha-beta pruning
 - ❖ Depth-limited search, iterative deepening
- ❖ Algorithm for random agent
 - ❖ Expectimax
- ❖ Algorithm for multi-agent search
 - ❖ Expectiminimax

Decision Theory and Game Theory

- ❖ Axiomatization of Expected Utility
 - ❖ Completeness, Transitivity, Independence, Continuity
 - ❖ Unicity of utility function up to positive affine transformation
 - ❖ Preference elicitation
- ❖ Game theory
 - ❖ Extensive form vs normal form
 - ❖ Best response, dominant/ dominated strategies
 - ❖ Nash equilibrium (pure or mixed) ✓
 - ❖ Pareto optimal, correlated equilibrium

Markov Decision Process

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model known
- ❖ Model
 - ❖ States, actions, transition function, reward function
- ❖ Algorithms
 - ❖ Policy evaluation ✓
 - ❖ Policy extraction ✓
 - ❖ Value iteration]
 - ❖ Policy iteration]

$$\pi : \mathcal{S} \rightarrow \mathcal{A}$$

$$v : \mathcal{S} \rightarrow \mathbb{R}$$

Reinforcement Learning

- ❖ Environment: single-agent, fully-observable state, stochastic transition, sequential, model unknown

- ❖ MDP Model, but unknown!

 - ❖ States, actions, transition function, reward function

- ❖ Algorithms

 - ❖ Policy evaluation with TD learning

 - ❖ Policy learning with Q-learning

 - ❖ Approximate Q-learning

$$\sum_{i=1}^n \overbrace{\phi_i(s,a)} \times w_i = \hat{Q}(s,a)$$

\uparrow

 - ❖ Action selection with ϵ -greedy or exploration function

Constraint Satisfaction

❖ CSP

- ❖ Set of variables, set of domains, set of constraints
- ❖ Find assignments to variables such that all constraints are satisfied

(x_1, \dots, x_n)

❖ Algorithms

- ❖ Backtracking search
 - ❖ Filtering, forward-checking, arc consistency, k-consistency
 - ❖ Ordering of variables and values
- ❖ Structure of constraint graph
 - ❖ Two-pass algorithm for tree-structured constraint graph
 - ❖ Cutset conditioning
- ❖ ~~Iterative~~ improvement

❖ Local search

Quiz: Search

- ❖ Consider a graph search problem where for every action, the cost is at least ϵ , with $\epsilon > 0$. Assume the used heuristic is consistent.
- ❖ Greedy graph search is guaranteed to return an optimal solution. **F**
- ❖ A^* graph search is guaranteed to return an optimal solution. **T**
- ❖ A^* graph search is guaranteed to expand no more nodes than depth-first graph search. **F**
- ❖ A^* graph search is guaranteed to expand no more nodes than uniform-cost graph search. **T**

$$g(n) + \downarrow h(n)$$

Quiz: A^* Heuristics

- ❖ Let H_1 and H_2 both be admissible heuristics.
 - ❖ $\max(H_1, H_2)$ is necessarily admissible ✗
 - ❖ $\min(H_1, H_2)$ is necessarily admissible T
 - ❖ $(H_1 + H_2)/2$ is necessarily admissible T
 - ❖ $\max(H_1, H_2)$ is necessarily consistent F

$$\forall n, h(n) \leq h^*(n)$$

$$\forall n, h(n) \leq c(n, n') + h(n')$$

Quiz: Search under Uncertainty

- ❖ You are given a game tree for which you are the maximizer, and in the nodes in which you don't get to make a decision an action is chosen uniformly at random amongst the available options. Your objective is to maximize the probability you win \$10 or more (rather than the usual objective to maximize your expected value).
- ❖ Running expectimax will result in finding the optimal strategy to maximize the probability of winning \$10 or more. F
- ❖ Running minimax, where chance nodes are considered minimizers, will result in finding the optimal strategy to maximize the probability of winning \$10 or more. F
- ❖ Running expectimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more. T
- ❖ Running minimax in a modified game tree where every pay-off of \$10 or more is given a value of 1, and every pay-off lower than \$10 is given a value of 0 will result in finding the optimal strategy to maximize the probability of winning \$10 or more. F

Quiz: Adversarial Search

- ❖ In the context of adversarial search, α - β pruning
 - ❖ can reduce computation time by pruning portions of the game tree T
 - ❖ is generally faster than minimax, but loses the guarantee of optimality F
 - ❖ always returns the same value as minimax for the root of the tree T
 - ❖ always returns the same value as minimax for all nodes of the tree F

Game Theory: Zero-Sum Game

- ❖ Two players choose simultaneously a coin of 10 cents, 50 cents or 1 dollar, which they show to each other.
- ❖ If they chose the same coin, player I wins. Otherwise, player II wins.
- ❖ Write this game in normal form. Is there any pure NE?
- ❖ Express a system of inequalities to find a mixed NE.

Handwritten notes and calculations for the game:

Normal Form Payoff Matrix:

	p_{II} 10	q_{II} 50	$(1-p_{II}-q_{II})$ 100
p_I 10	10 ^x	-10 ^o	-10 ^o
q_I 50	-50 ^o	50 ^x	-50 ^o
$(1-p_I-q_I)$ 100	-100 ^o	-100 ^o	100 ^x

System of Inequalities for Mixed NE:

$$\begin{cases} 10p - 10q - 10(1-p-q) \geq v \\ -50p + 50q - 50(1-p-q) \geq v \\ -100p - 100q + 100(1-p-q) \geq v \\ 0 \leq p \leq 1, 0 \leq q \leq 1 \end{cases}$$

Further Elimination:

$$\begin{array}{l|l} 0 \leq p & p \leq 1 \\ \frac{1}{20}v + \frac{1}{2} \leq p & p \leq -\frac{1}{20}v - q + \frac{1}{2} \\ 0 \leq q \leq 1 & q \geq \frac{1}{100}v + \frac{1}{2} \end{array}$$

Final Inequalities:

$$\begin{cases} 20p \geq v + 10 \\ 100q \geq v + 50 \\ -200p \geq v + 200q - 100 \end{cases}$$

Value of game: v

$$\max\left(0, \frac{1}{20}v + \frac{1}{2}\right) \leq p \leq \min\left(1, -\frac{1}{200}v - 9 + \frac{1}{2}\right)$$

$$\left\{ \begin{array}{l} 0 \leq -\frac{1}{200}v - 9 + \frac{1}{2} \\ \frac{1}{20}v + \frac{1}{2} \leq 1 \\ \frac{1}{20}v + \frac{1}{2} \leq -\frac{1}{200}v - 9 + \frac{1}{2} \\ 0 \leq p \leq 1 \\ p \geq \frac{1}{100}v + \frac{1}{2} \end{array} \right.$$

$$\left\{ \begin{array}{l} 0 \leq p \\ \frac{1}{100}v + \frac{1}{2} \leq p \\ p \leq -\frac{1}{200}v + \frac{1}{2} \\ p \leq -\frac{11}{200}v \\ p \leq 1 \end{array} \right.$$

$$\frac{1}{20}v + \frac{1}{2} \leq 1 \quad v \leq -100 \quad \downarrow$$

$$\begin{array}{l} 0 \leq -\frac{1}{200}v + \frac{1}{2} \quad v \leq 100 \\ 0 \leq -\frac{11}{200}v \quad v \leq 0 \\ \frac{1}{100}v + \frac{1}{2} \leq -\frac{1}{200}v + \frac{1}{2} \\ \cancel{\frac{1}{20}v} \leq 0 \end{array}$$

$$\begin{array}{l} \frac{1}{100}v + \frac{1}{2} \leq -\frac{11}{200}v \quad \frac{13}{200}v \leq -\frac{1}{2} \\ \frac{1}{100}v + \frac{1}{2} \leq 1 \quad v \leq 50 \\ \frac{1}{20}v + \frac{1}{2} \leq 1 \quad v \leq 10 \end{array}$$

$$\max_{\sigma} \min_{\sigma'} \sum_{ij} \overset{\text{probabilities}}{\sigma_i \sigma'_j} u_{ij} = v$$

Quiz: MDP

❖ For Markov Decisions Processes (MDPs), we have that:


- T ❖ A small discount (close to 0) encourages shortsighted, greedy behavior.
- T ❖ A large, negative living reward ($\ll 0$) encourages shortsighted, greedy behavior.
- F ❖ A negative living reward can always be expressed using a discount < 1 .
- F ❖ A discount < 1 can always be expressed as a negative living reward.

Quiz: MDP

- ❖ Value iteration can converge only if the discount factor (γ) satisfies $0 < \gamma < 1$.
F
- ❖ Policies found by value iteration may be superior to policies found by policy iteration.
F
- ❖ Policies found by policy iteration may be superior to policies found by value iteration.
F
- ❖ In some problems, value iteration can converge even though policy iteration may not. $\gamma = 1$
T

Quiz: Reinforcement Learning

- ❖ Assume that the agent observes the true reward with some
- ✓ Gaussian noise $\mathcal{N}(0,1)$, Q-learning would still converge
- ❖ Q-learning can learn the optimal Q-function Q^* without
- ✓ *off-policy* ever executing the optimal policy.]
- ❖ If an MDP has a transition model T that assigns non-zero
- ✗ probability for all triples $T(s, a, s')$ then Q-learning will fail.
- ❖ In Q-learning, we decide to explore every k steps, i.e., if $t = 0 [k]$ we choose a random action with a uniform
- ✗ distribution, otherwise we choose the greedy action. This version would still converge.




Quiz: CSP

- ❖ Assume given a CSP whose constraint graph is given below and that all the variables have the same domain.
- ❖ What is the complexity of solving it with a direct application of backtracking search? $O(d^{29})$
- ❖ Which efficient strategy could you apply to solve it? What would be the complexity?

29 variables

d size of domain

$O(d \times d^2 \times \cancel{28})$



CSP Problem: Job Scheduling

❖ When can I move in?

Task	Description	Duration	Predecessor
a	Erecting walls	7	none
b	Carpentry for roof	3	a
c	Roof	1	b
d	Installations	8	a
e	Facade painting	2	c & d
f	Windows	1	c & d
g	Garden	1	c & d
h	Ceilings	3	a
i	Painting	2	f & h
j	Moving in	1	i

